
COL 331: OS ASSIGNMENT 3

ASSIGNMENT 3: REPORT

2023-04-30

REEDAM DHAKE — 2020CS10372

NIKHIL UNAVEKAR — 2020CS10363

Contents

1	Changes Made	2
2	Address Space Layout Randomization	2
3	Buffer Overflow Attack in XV6	6

1 Changes Made

- Added 3 function `read_aslr_flag`, `set_randoffset_text` and `set_randoffset_stack` in `exec.c`
- Added 2 functions `get_pa` and `new_virtual_memory_loader` in `vm.c`
- Defined 3 files - `alsr_flag` (for ALSR value), `gen_exploit.py` and `random.h`. `gen_exploit.py` generates payload and `random.h` is the random integer generator.

2 Address Space Layout Randomization

CHANGES IN VM.C

```
1  uint get_pa(pde_t *pte, pde_t *pgdir, char *first_va, int alloc){
2      if((pte = walkpgdir(pgdir, (char *)first_va, 0)) == 0)
3          panic("loaduvm: address should exist");
4      else
5          return PTE_ADDR(*pte);
6  }
7
8  int new_virtual_memory_loader(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz){
9      uint i, pa, n, va=(uint)addr;
10     pte_t *pte;
11     uint first_va = PGROUNDDOWN(va);
12     uint pg_offset = va - first_va;
13     pa = get_pa(pte, pgdir, (char*)first_va, 0);
14     n = (sz < PGSIZE - pg_offset) ? sz : PGSIZE - pg_offset;
15     if(readi(ip, P2V(pa + pg_offset), offset, n) != n)
16         return -1;
17     offset += n;
18     sz -= n;
19     for(i = PGSIZE; sz > 0; i += PGSIZE){
20         pa = get_pa(pte, pgdir, (char*)(first_va + i), 0);
21         n = (sz < PGSIZE) ? sz : PGSIZE;
22         if(readi(ip, P2V(pa), offset, n) != n)
23             return -1;
24         offset += n;
25         sz -= n;
26     }
27     return 0;
28 }
```

- The `get_pa` function is called by the `new_virtual_memory_loader` function to retrieve the physical address corresponding to a given virtual address. It takes as input a page table entry pointer `pte`, a pointer to the first byte of the page containing the virtual address `first_va`, a page directory entry pointer `pgdir`, and an integer `alloc` for new page allocation.
- It first calls the `walkpgdir` function to obtain a pointer to the page table entry corresponding to the virtual address `first_va` in the page directory `pgdir`. If `walkpgdir` returns a null pointer, indicating that the page table entry is not present, error is returned. Else, the function returns the physical address corresponding to the page table entry by making the page table entry with the `PTE_ADDR`, which extracts the physical address from the page table.

- In the `new_virtual_memory_loader` function, the `thiss` function is used to retrieve the physical address of the page containing the first word of the virtual address `addr`, followed by retrieving the physical addresses of each subsequent page that needs to be loaded with data from the file.

`new_virtual_memory_loader` takes a page directory entry pointer `pgdir`, a virtual address `addr`, an inode `ip`, an offset `offset`, and a size `sz` as inputs.

- This function loads `sz` bytes of data from the file referred to by the inode `ip` at the specified offset into the virtual memory at the address `addr`, using the page directory `pgdir`.
- It first calculates the page offset `pg_offset` of the virtual address `addr` from the beginning of the page, and then retrieves the physical address `pa` corresponding to the page containing the first byte of `addr` using the `get_pa` function as described above.
- It then reads the first `n` bytes of data from the file at the specified offset into the physical memory starting at `P2V(pa + pg_offset)`, where `P2V` is used to convert the a physical address to a virtual address. It then subtracts `n` from `sz` to indicate that `n` bytes have been loaded.
- Finally it loops to load the remaining bytes. In each iteration, the function calculates the physical address `pa` of the next page of virtual memory by adding the page size `PGSIZE` to the previous physical address `pa`. Then it reads the next `n` bytes of data from the file at the specified offset into the physical memory starting at `P2V(pa)`. The function then subtracts `n` from `sz` to indicate that `n` more bytes have been loaded. The loop continues until all `sz` bytes have been loaded. Error returns -1 else return 0.

CHANGES IN EXEC.C

```

1  int read_aslr_flag(struct inode *ip){
2      char temp[1]={0};
3      ip = namei("aslr_flag");
4      ilock(ip);
5      readi(ip,&temp[0],0,8);
6      iunlockput(ip);
7      if(temp[0]=='1'){
8          return 1;
9      }
10     else{
11         return 0;
12     }
13 }
14
15 int set_randoffset_text(int asfl){
16     int offset1=random_num_gen(1,4000);
17     offset1*=32;
18     if(asfl==0){
19         return 0;
20     }
21     return offset1;
22 }
23
24 int set_randoffset_stack(int asfl){
25     int extra_pgs=random_num_gen(2,30);
26     int offset2=2+extra_pgs;
27     if(asfl==0){
28         offset2=2;

```

```

29     }
30     return offset2;
31 }

1  int
2  exec(char *path, char **argv)
3  {
4      ...
5      begin_op();
6      int asfl=read_aslr_flag(ip);
7      ....
8      // Check ELF header
9      int offset1=set_randoffset_text(asfl);
10     sz = 0;
11     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
12         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
13             goto bad;
14         if(ph.type != ELF_PROG_LOAD)
15             continue;
16         if(ph.memsz < ph.filesz)
17             goto bad;
18         if(ph.vaddr + ph.memsz < ph.vaddr)
19             goto bad;
20         if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz+offset1)) == 0)
21             goto bad;
22         if(ph.vaddr % PGSIZE != 0)
23             goto bad;
24         if(new_virtual_memory_loader(pgdir, (char*)(ph.vaddr+offset1), ip, ph.off, ph.filesz) < 0)
25             goto bad;
26     }
27     iunlockput(ip);
28     end_op();
29     ip = 0;
30     int offset2=set_randoffset_stack(asfl);
31     sz = PGROUNDUP(sz);
32     if((sz = allocuvm(pgdir, sz, sz + offset2*PGSIZE)) == 0)
33         goto bad;
34     clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
35     sp = sz;
36     ....
37 }

```

- Start defines the 3 functions which are explained below and which we will be using the core exec.c function.
- The read_aslr_flag reads the value of a file named "aslr_flag" and the flag value, 1 indicates that Address Space Layout Randomization (ASLR) is enabled, and 0 means its disabled. It takes a pointer to an inode structure as an argument, which is used to access the file.
- The set_randoffset_text generates a random offset for the program text section based on the value of asfl. If asfl is 0, it returns 0, otherwise, it generates a random number between 1 and 4000, multiplies it by 32, and returns it.
- The set_randoffset_stack generates a random offset for the stack section based on the value of flag asfl. If asfl is 0, it returns 2, otherwise, it generates a random number between 2 and 30 and adds 2 to it before returning the result. The exec function loads and executes a new program by replacing the

current process's memory image with the program that is to be executed. This function takes the path of the executable and its arguments as parameters.

- First, it reads the value of the ASLR flag by calling the `read_aslr_flag` function, which opens and reads a file named "aslr_flag" to check if the ASLR feature is enabled or not. If ASLR is disabled, the offset values for text and stack segments will be set to their default values.
- After obtaining the ASLR flag, `exec` opens the executable file by calling `namei` function and checks the ELF header of the program that is to run. If the ELF header is invalid, the function returns an error. Next it sets up a new kernel page table by calling `setupkvm` function. The old page directory is saved to `oldpgdir` to be restored later.
- Then, `exec` calculates a random offset value for the text segment by calling the `set_randoffset_text` function, which generates a random number between 1 and 4000 and multiplies it by 32. If ASLR is disabled, offset will be set to 0.
- If an attacker attempts to launch a buffer overflow attack against a program with ASLR enabled, they will not be able to predict the address where the malicious code will be loaded, as the address is randomized every time this program is executed.

The random integer generator is described below -

RANDOM.H

```
1  int rand1(int rr){
2      int rand = rr;
3      rand=(rand+(ticks))%(85663);
4      rand *= 12;
5      rand += 8965;
6      rand ^= 1463588;
7      return rand;
8  }
9
10 int random_num_gen(int mini,int maxi){
11     int rand = rand1(2);
12     for(int i=0;i<ticks*5;i++){
13         rand=rand1(rand);
14     }
15     rand = rand % 1000000009;
16     rand=rand%(maxi-mini+1);
17     printf("%d\n",rand);
18     return mini+rand;
19 }
```

The `rand1` function takes an integer `rr` as input and uses it as the seed for the random number generator. It generates a random number by performing a series of arithmetic operations on the seed value, including adding the current value of the ticks of system, adding a constant value, performing a bitwise XOR with another constant value, etc. The resulting value is returned as the random number. This function is called by `random_num_gen` function. It first calls `rand1` with a seed value of 2, and then iteratively calls `rand1` on the result of the previous call `ticks*5` number of times to generate a more random value. It then takes the modulus of the resulting value with a large number. Finally, it takes the modulus of the range between the given minimum and maximum values and adds the minimum value to the result to generate a random number within the desired range of max and min.

3 Buffer Overflow Attack in XV6

```
1 import sys
2 f = open("payload", "wb")
3 f.write(b"aaaaaaaaaaaa" + b'a'*(int(sys.argv[1])) + b"\x00\x00" * 2)
4 f.close()
```

- The above script can be used for to generate a payload for buffer overflow attack . As we know in a buffer overflow attack, an attacker tries to overwrite a buffer with more data than the buffer can hold, with the goal of overwriting adjacent memory locations. b"aaaaaaaaaaaa" string of 12 bytes (ASCII characters 'a') serves as a buffer overflow trigger.
- The length of this string is the minimum amount of data required to overflow the buffer.
- Then it creates a string of variable length, depending on the command line argument passed to the script. This string is the payload that will overflow the buffer. It consists of 'a' characters repeated input times. sys.argv[1] is the first argument. b"\x00\x00" * 2 adds two null bytes to the end of the payload.
- Null bytes can be used to terminate strings, so adding them can help ensure that the payload does not corrupt other data beyond the intended buffer. An attacker could generate a payload of arbitrary length, and then pass the length as a command line argument when running the actual script.
- The resulting payload file could then be used as input to a vulnerable program that is to receive a buffer overflow attack. When the program reads in data from the payload file, it will eventually hit the b"aaaaaaaaaaaa" string, which will trigger the overflow and allow the payload to overwrite adjacent memory locations.